

## MACM 316 – Computing Assignment 6

- Read the *Guidelines for Assignments* first.
- Submit a one-page PDF report to Canvas and upload your Matlab scripts (as m-files). Do not use any other file formats.
- Keep in mind that Canvas discussions are open forums.
- You must acknowledge any collaborations/assistance from colleagues, TAs, instructors etc.

### Machine learning

In this assignment, you will use least-squares fitting to teach your computer to distinguish between red and blue points in 2D. This is an example of linear discriminant analysis.

First, download the file *dataset.mat*. This contains two tall matrices, *training\_set* and *test\_set*. Your goal is to train your computer using the training set and then use the test set to see how well it has learned.

### Description of the dataset

Each dataset has three columns. The training set has 2000 rows (i.e. 2000 red and blue points) and the test set has 400 points. Let

```
1 X = training_set(:, [1 2]); y = training_set(:, 3);
```

Each row of  $\mathbf{X}$  is the  $(x, y)$  coordinates of a point in 2D. The corresponding entry in  $\mathbf{y}$  is either 0 or 1, with 1 representing the colour blue and 0 representing the colour red.

Use Matlab's *load* command to import the datasets. Next, assemble  $\mathbf{X}$  and  $\mathbf{y}$  as above and use the *plot* command to reproduce the following plot:




Figure 1: training dataset

## Training phase

Define the matrix  $\mathbf{A}$  as follows:

```
1 A = [ ones(2000,1) X ];
```

Note that  $\mathbf{A}$  is a  $2000 \times 3$  matrix and  $\mathbf{y}$  is a  $2000 \times 1$  vector. The objective of the training is to find a vector  $\boldsymbol{\beta}$  such that  $\mathbf{A}\boldsymbol{\beta} \approx \mathbf{y}$ . Using least-squares fitting (implemented via backslash), compute such a vector, and call it  $\hat{\boldsymbol{\beta}}$ . Use this to compute the RSS (Residual Sum of Squares) value

$$\text{RSS} = \|\mathbf{y} - \mathbf{A}\hat{\boldsymbol{\beta}}\|^2.$$

**Record this value in your report.** This gives an idea of how good  $\hat{\boldsymbol{\beta}}$  is at predicting a value in  $\mathbf{y}$  given the corresponding row in  $\mathbf{A}$ .

Next, in Figure 1 add a black line given by the following equation

$$\hat{\beta}_1 + \hat{\beta}_2 x_1 + \hat{\beta}_3 x_2 = 1/2. \quad (1)$$

**What is this line saying about the red and blue points? Explain your answer, and include both the figure and your explanation in the report.**

## Testing phase

It is now time to see how well the computer has learned. Let

```
1 z = test_set(:,3);
2 B = [ ones(400,1) test_set(:,[1 2]) ];
```

Note that  $\mathbf{z}$  is a  $400 \times 1$  vector and  $\mathbf{B}$  is a  $400 \times 3$  matrix similar to  $\mathbf{A}$ . Write a few lines of code to compute the vector  $\mathbf{v} = \mathbf{B}\hat{\boldsymbol{\beta}}$  and the vector  $\mathbf{z}$  defined by

$$\hat{z}_j = \begin{cases} +1 & \text{if } v_j \geq 1/2 \\ 0 & \text{if } v_j < 1/2 \end{cases}.$$

Compute the error of the prediction as follows:

$$\text{Err} = \frac{1}{400} \sum_{j=1}^{400} |z_j - \hat{z}_j|.$$

**Record this value in your report, and discuss. Also, generate a plot of the test dataset where the points are labelled according to the prediction  $\hat{z}$  along with the discriminant line defined in equation (1) above.**

## Bonus! (2 additional marks)

You will hopefully have seen that linear discriminant analysis can do quite well for training a computer to distinguish certain datasets. Indeed, it is a powerful technique. However, it does not work well for all datasets. Devise a dataset for which you expect it would not work so well, and run your code on this dataset to demonstrate this observation.